



Need to Know Microsoft 365 Webinar September 2026

@directorcia

<http://about.me/ciaops>

Web cast has started

Web cast is being recorded. Note recording may be made public

If you can't hear anything check your speaker settings

For questions after the event:

Email : director@ciaops.com

X : [@directorcia](#)



directorcia

Your email address

Subscribe

X

CIAOPS

Products

+

Copilot

1-6 of 6 products

Sort by

>

Tags

>

Contains

>

Price

>

Copilot Notebooks:
A Beginner's Guide

Introduction to Copilot Notebooks

A\$9+

AI REVENUE & COPILOT READINESS BUNDLE

MSP AI Playbooks

A\$18+

AI as a Revenue Stream for SMB-Focused MSPs:
A Step-by-Step Playbook

AI as a Revenue Stream for MSPs

A\$9+

Your Copilot ROADMAP

Your Copilot Roadmap

A\$9+

How to Effectively use Microsoft 365 Copilot

How to Effectively Use Microsoft 365 Copilot

A\$9+

COPILOT READINESS ASSESSMENT FOR SMB

Copilot Readiness Assesment

A\$9+

+

Security

1-6 of 6 products

Sort by

>

Tags

>

Contains

>

Price

>

SMB Identities in M365 Business Premium from Token Theft

Protecting SMB Identities in M365 Business Premium from Token Theft

A\$9+

Cybersecurity Verification Playbook

Cybersecurity Verification Playbook

A\$9+

MICROSOFT 365 BUSINESS PREMIUM COMPLIANCE GUIDE FOR AUSTRALIAN SMALL BUSINESSES
by Robert Crane

Microsoft 365 Business Premium Compliance Guide

A\$9+

Aligning 2025-26

Aligning 2025-26

A\$9+

Implementing ACSC

Implementing ACSC

A\$9+

RESPONSE PLAN

Response Plan

A\$9+

<https://blog.ciaops.com/2022/07/29/join-my-teams-shared-channel/>

Join my
free
shared
Channel



Agenda

- Microsoft 365 Update
- Beyond Copilot Chat
- Q & A

Microsoft 365 Update



News

- Introducing the new Copilot with Home, Code and Autopilot
 - <https://blogs.microsoft.com/blog/2026/09/25/introducing-the-new-copilot-with-home-code-and-autopilot/>
- Available today: OpenAI's GPT-6.1 Sol and Claude Sonnet 5.5 in Microsoft Copilot
 - <https://techcommunity.microsoft.com/blog/microsoft-copilot-blog/available-today-openais-gpt-6-1-sol-and-claude-sonnet-5-5-in-microsoft-copilot/4560801>
- EWS Deprecation Is Here – What This Means To You
 - <https://techcommunity.microsoft.com/blog/exchange/ews-deprecation-is-here-%E2%80%93-what-this-means-to-you/4561431>
- Insights from the 2026 Microsoft Digital Defense Report
 - <https://www.microsoft.com/en-us/security/blog/2026/10/01/insights-from-the-2026-microsoft-digital-defense-report/>
- Phishing Abuses RMM Tools for Persistent Access
 - <https://www.microsoft.com/en-us/security/blog/2026/09/29/phishing-abuses-rmm-tools-persistent-access/>

On prem AI



What Happens in Foundry Local

```
</> foundry model run phi-4-mini
```

FOUNDRY LOCAL THEN



Loads the model



Places model weights into memory



Sends computations to the best available silicon

COMPUTE TARGET ORDER

GPU

Preferred

NPU

Supported

CPU

Fallback

Microsoft's design leverages local AI acceleration wherever it's available — GPU first, NPU next, CPU only as a fallback.

Install Foundry Local & Run a Model

From a clean, GPU-ready machine to a model answering prompts — in five commands.

STEP BY STEP

1

Install Foundry Local

```
winget install Microsoft.FoundryLocal
```

2

Confirm the service is up

```
foundry service status
```

3

Browse available models

```
foundry model list
```

4

Download and run a model

```
foundry model run phi-3.5-mini
```

5

Chat, or point Open WebUI at it

```
http://localhost:5273/v1
```

BEFORE YOU START

- Windows 11 with GPU drivers installed
- nvidia-smi returns a healthy GPU
- ~10 GB free disk for model weights

CONFIRM IT'S ON THE GPU

```
nvidia-smi -l 1
```

VRAM usage should jump when the model loads. If it doesn't, Foundry Local fell back to CPU.

Start small: a 3B model proves the pipeline before you commit VRAM to anything larger.

Where Open WebUI Fits

Open WebUI is not the AI model — it is the experience layer on top of your models.

THE FOUNDRY LOCAL STACK

You

Asks the question

Open WebUI

Chat interface / UX layer

Foundry Local

Runtime that serves the models

Phi / Qwen / DeepSeek

The models themselves

GPU

Does the actual compute

THE GPU

Does the work

Raw parallel compute turns tokens into throughput.

THE MODEL

Generates the answer

Phi, Qwen or DeepSeek produce the actual response.

OPEN WEBUI

Makes it usable

Chat, history, model switching, file uploads — no CLI required.

Without the UX layer, the capability exists — but nobody uses it.

Architecture Overview

The full stack you end up with, from the operating system down to the GPU.

THE STACK, TOP TO BOTTOM

Windows 11

Host operating system

Docker Desktop

Container runtime

Open WebUI Container

Chat interface

Foundry Local

Inference engine

Phi / Qwen / DeepSeek

The models

RTX GPU

Compute hardware

OPEN WEBUI

The chat interface

Everything the user touches: conversations, history, uploads, model selection.

FOUNDRY LOCAL

The inference engine

Loads models onto the GPU and serves an OpenAI-compatible API on localhost.

The container talks to Foundry Local on the host via `host.docker.internal:5273`.

One machine, six layers — interface on top, silicon at the bottom.

DEMO

Install Ollama

Download the Windows installer, accept the defaults, then verify from PowerShell.

STEP BY STEP

1

Download the installer

ollama.com/download — Ollama Windows Installer

2

Run OllamaSetup.exe

Launch it from your Downloads folder

3

Accept the defaults

No custom options are needed

4

Open a new PowerShell window

A new window picks up the updated PATH

5

Check the installed version

Confirms Ollama is installed and on your PATH

RUN THE INSTALLER

 **OllamaSetup.exe**

Accept every default prompt as it appears.

VERIFY FROM POWERSHELL

ollama --version

Expected output:

ollama version 0.x.x

No version number? Reopen PowerShell, then re-run the installer if it still fails.

A version number means Ollama is installed and ready to pull your first model.

DEMO

Local AI vs Microsoft 365

A local AI model has no native understanding of your Microsoft 365 tenant.

WHERE LOCAL AI FALLS SHORT

- Poor access to Microsoft 365 data
- Minimal audit capability
- No native Microsoft Graph grounding
- Complex authentication management
- Limited permission trimming
- Significant operational overhead
- Weak governance and compliance
- Reduced user experience
- No Purview integration
- Difficulty scaling across an organisation

WHY THE GAP EXISTS

Microsoft 365 Copilot is deeply integrated with Microsoft Graph, Entra ID, SharePoint permissions, Teams, Outlook and Purview. A local AI must be explicitly connected to every one of those services.

THE TRADE-OFF

Privacy and offline control, minus tenant awareness.

Connecting a local model to M365 data is integration work, not a setting.

Local AI is strong offline - tenant-aware work still needs Microsoft 365 integration.

The Identity Problem

With Microsoft 365 Copilot, the identity chain is already in place before the model is asked anything.

THE COPILOT IDENTITY CHAIN

User



Entra ID



Microsoft Graph



SharePoint

Teams

Outlook

OneDrive

WHAT THAT MEANS

- The identity is already known.
- The permissions are already known.
- The security context already exists.
- Everything already trusts everything else.

A local AI starts with none of this - identity has to be built from scratch.

Copilot inherits trust; local AI has to earn it, service by service.

Takeaways

- Local AI runs on your own hardware - private and offline
- Foundry Local and Ollama make running models easy
- GPU first, NPU next, CPU fallback - silicon drives speed
- Open WebUI is the layer that makes local models usable
- Local models have no awareness of your M365 tenant
- Copilot inherits identity; local AI must earn it
- Local for privacy, Microsoft 365 for tenant-aware work

Resources

- Foundry Local documentation - <https://learn.microsoft.com/en-us/azure/foundry-local/>
- Get started with Foundry Local - <https://learn.microsoft.com/en-us/azure/foundry-local/get-started>
- Foundry Local on Windows AI - <https://learn.microsoft.com/en-us/windows/ai/foundry-local/get-started>
- Foundry Local CLI reference - <https://learn.microsoft.com/en-us/azure/foundry-local/reference/reference-cli>
- Open WebUI documentation - <https://docs.openwebui.com/>
- Ollama - <https://ollama.com/>
- Microsoft 365 Copilot overview - <https://learn.microsoft.com/en-us/microsoft-365/copilot/microsoft-365-copilot-overview>

CIAOPS Resources



- Blog – <http://blog.ciaops.com>
- Free Office 365, Azure video tutorials – <http://www.youtube.com/directorciaops>
- Free documents, presentations, eBooks – <https://github.com/directorcia/general>
- Office 365, Azure, Cloud podcast – <http://ciaops.podbean.com>
- Publications– <http://directorcia.gumroad.com>
- CIAOPS Patron Community – <http://www.ciaopspatron.com/>
- CIAOPS AI Dojo – <https://www.ciaops.com/ai-dojo>
- LinkedIn - <http://www.linkedin.com/in/ciaops>
- Github - <https://github.com/directorcia/>



[Twitter](#)
[@directorcia](#)

[Facebook](#)
<https://www.facebook.com/ciaops>

[Email](#)
director@ciaops.com

[Teams](#)
admin@ciaops365.com



Get access to the latest
information by becoming a
Patron

<http://www.ciaopspatron.com>



Questions

That's all folks!

Thanks for attending